Classification

A task to predict (value) classes

- Digits
- Faces
- Nodes
- Fruits

Binary Classifier:

Task to predict if a given input is a class (1) or not (0)

Basically an if statement on a class so, predict() would give true or false

- Performance measures

Cross validation
(oh boy)

Split the data into multiple folds and then train the model on every fold except one, repeat this ~~etc~~ for all folds.

Then you will have an array of accuracys for each fold.

ex: [0.91, 0.88, 0.99]

Pseudo-code - Python

Given K

new_Data = Data / K # Split into folds

for each fold in new_Data:

~~EEE: or~~ # Split training and testing
* data-train = new_Data - fold
* data-test = fold

acc.append(Model.fit(d*train, d*test).accuracy)

print(acc) # [0.91, 0.88, 0.99]

Im(~~Cross validation~~) Classification acc. is not the preferd performance measure, specially with skewed datasets (some classes are much more frequent than others)

In some cases when a class occupies 90% of the data, the model will just always guess that class and get an acc. of 90%.