


What is Reinforcement Learning

- The idea behind RL is that an agent (an AI) will learn from the environment by interacting with it (through trial and error) and receiving rewards (negative or positive) as feedback for performing actions.

How does it work



- Our Agent receives state S_t from the environment (we receive the first frame of our game (environment))
- Based on that state S_t , the agent takes action A_t (Agent will move to the right)
- Environment goes to a new state S_{t+1} (new frame)
- Environment gives some reward R_t to the agent (we're not dead (Positive reward))

- The agent's goal is to maximize its cumulative reward, called expected return