ϵ-greedy policy

An effective way to balance exploration and exploitation in RL

It is based on the idea that an agent should sometimes take an exploratory action to learn more.

The ϵ in ϵ-greedy stands for "epsilon", which represents a small probability value between 0 and 1.

A ϵ value of 0.1, the agent will take the best action 10% of the time.

- With $1-\epsilon$ probability - we do exploitation
- With ϵ probability - we do exploration

Adjusted ϵ over time function: