

Off-policy vs On-policy

- Off-policy: using a different policy for acting and training.

For instance we can use ϵ -greedy policy for training and a greedy policy (that always selects best action) for acting.

- On-policy: Using the same policy for acting and training.

Glossary — // — // — // — //

Explor - greedy strategy

- Common exploration strategy used reinforcement learning that involves balancing exploration and exploitation.
- Chooses the action with the highest expected reward with a probability of $1 - \epsilon$.
- Chooses a random action with a probability of ϵ .
- ϵ is typically decreased over time to shift focus towards exploitation.

Greedy Strategy

- Always chooses exploitation.
- Does not include exploration.
- Can be disastrous in environments with uncertainty.

— // — // — // — //